



Kubernetes Workloads API: ReplicaSet, Deployment StatefulSet, DaemonSet, Job, CronJob

Check GitHub for helpful DevOps tools:

Michael Robotics


Hi, I'm Michal. I'm a Robotics Engineer and DevOps enthusiast. My mission is to create skill-learning platform that combats information overload by adhering to the set of principles: simplify, prioritize, and execute.

 <https://github.com/MichaelRobotics>



Ask Personal AI Document assistant to learn
interactively (FASTER)!

- 1 Download PDF
 - 2 Go to website
 - 3 Browse file
 - 4 Chat with Document
- 1 <https://github.com/MichaelRobotics/DevOpsTools/blob/main/KubernetesWorkloads.pdf>
- 2 | Click there to go to ChatPdf website
- 3
- 4 Ask questions about document!



Completly new to Linux and Networking?

Essential for this PDF is a thorough knowledge of networking. I highly recommend the HTB platform's networking module, which offers extensive information to help build a comprehensive understanding.

HTB - Your Cyber Performance Center

We provide a human-first platform creating and maintaining high performing cybersecurity individuals and organizations.

 <https://www.hackthebox.com/>



What is Kubernetes?

Kubernetes is an open-source platform that automates the deployment, scaling, and management of containerized applications. It helps manage clusters of nodes running containers, ensuring efficient and reliable operation.

How Kubernetes clusters are made?

Kubernetes clusters consist of a control plane and multiple worker nodes. The control plane manages cluster operations, while worker nodes run the actual container workloads.

Why and When use Kubernetes

Kubernetes is ideal for deploying scalable, resilient, and automated containerized applications. It is used when managing multiple containers across different environments is necessary.

Example: Running a microservices-based e-commerce platform that scales up during peak hours.

System Requirements

- RAM: 2 GB per node (1 GB can work for testing but may lead to limited performance)
- 10 GB free storage
- Ubuntu

Kubernetes: Main components & packages

- **kube-apiserver:** Central management component that exposes the Kubernetes API; acts as the front-end for the cluster.
- **etcd:** Distributed key-value store for storing all cluster data, ensuring data consistency across nodes.
- **kube-scheduler:** Assigns pods to available nodes based on resource requirements and policies.
- **kube-controller-manager:** Manages core controllers that handle various functions like node status, replication, and endpoints.
- **kubelet:** Agent that runs on each node, responsible for managing pods and their containers.
- **kube-proxy:** Manages networking on each node, ensuring communication between pods and services within the cluster.

Kubernetes Workloads API: Intro

1) What are Kubernetes Workloads

The Kubernetes Workload API is a key component of Kubernetes, focusing on managing and orchestrating workloads (applications and services) that run on a Kubernetes cluster.

2) k8s Workloads

Pod - The smallest deployable unit in Kubernetes, representing one or more containers that share the same network namespace and storage volumes.

ReplicaSet - Ensures a specified number of identical Pod replicas are running at any given time.

Deployment - Manages stateless applications and provides declarative updates for Pods and ReplicaSets.

StatefulSet - Manages stateful applications where each instance (Pod) requires a unique identity or stable storage.

DaemonSet - Ensures that a copy of a Pod runs on all (or some) nodes in the cluster.

Job - Manages batch or one-off tasks that need to complete successfully.

CronJob - Extends Job by running tasks on a schedule.

Kubernetes Workload API: ReplicaSet

Download project repo

```
git clone https://github.com/vfarcic/kubernetes-demo
```

As covered in intro, ReplicaSet is responsible to assure that specified number of pods runs all the time.

apply

```
cd kubernetes-demo
kubectl apply --filename replicaset/base.yaml
```

Install krew and tree:

```
curl -O https://raw.githubusercontent.com/MichaelRobotics/Kubernetes/main/Tools/InstallTree
bash InstallTree
source ~/.bashrc
```

Check ownership of resource. It's easy to find out, that replicaset is responsible for creating pods (number of those pods is specified in yaml file)

```
kubectl tree replicaset silly-demo
```

```
controlplane $ kubectl tree replicaset silly-demo
NAMESPACE  NAME                                READY  REASON  AGE
default    ReplicaSet/silly-demo             -      -        49s
default    Pod/silly-demo-dvbr5              True   -        49s
default    Pod/silly-demo-l62bc              True   -        49s
default    Pod/silly-demo-xbks6              True   -        49s
controlplane $
```

Delete one of the pods and check what will happen:

```
kubectl delete pod <pod_name>
```

```
controlplane $ kubectl tree replicaset silly-demo
NAMESPACE NAME READY REASON AGE
default   ReplicaSet/silly-demo - 8m7s
default   Pod/silly-demo-9t6mf True 3s
default   Pod/silly-demo-162bc True 8m7s
default   Pod/silly-demo-xbks6 True 8m7s
controlplane $
```

A new Pod was created, and the total count remains stable. Now, update the number of Pods by applying a ReplicaSet YAML with the same name but a different replica count:

```
kubectl apply --filename replicaset/replicas.yaml
```

Check any changes

```
kubectl tree replicaset silly-demo
```

```
controlplane $ kubectl tree replicaset silly-demo
NAMESPACE NAME READY REASON AGE
default   ReplicaSet/silly-demo - 95s
default   Pod/silly-demo-9qntc True 95s
default   Pod/silly-demo-jq9js True 32s
default   Pod/silly-demo-qcqhF True 95s
default   Pod/silly-demo-w44ft True 22s
default   Pod/silly-demo-zq7rn True 22s
controlplane $ kubectl apply --filename replicaset/replicas.yaml
replicaset.apps/silly-demo unchanged
```

New pods were created, reflecting the replica count change. Next, update the image type:

```
kubectl apply --filename replicaset/image.yaml
```

```
kubectl tree replicaset silly-demo
```

```
controlplane $ kubectl tree replicaset silly-demo
NAMESPACE NAME READY REASON AGE
default   ReplicaSet/silly-demo - 5m55s
default   Pod/silly-demo-9qntc True 5m55s
default   Pod/silly-demo-jq9js True 4m52s
default   Pod/silly-demo-qcqhF True 5m55s
default   Pod/silly-demo-w44ft True 4m42s
default   Pod/silly-demo-zq7rn True 4m42s
```

Pod images remain unchanged as ReplicaSet only manages pod count in real-time. To use a different image, delete all existing pods.

Kubernetes Workload API: Deployment

delete previous ReplicaSet

```
kubectl delete --filename replicaset/image.yaml
```

Deployment manages rolling updates, scaling, and rollback for stateless applications

apply

```
kubectl apply --filename deployment/base.yaml
```

As you can see, Deployment make use of ReplicaSet to manage pods

```
kubectl tree replicaset silly-demo
```

```
controlplane $ kubectl tree deployment silly-demo
NAMESPACE NAME                                READY REASON AGE
default   Deployment/silly-demo                    -      -      66s
default   └─ ReplicaSet/silly-demo-5b764b57cc      -      -      66s
default   │   └─ Pod/silly-demo-5b764b57cc-gxtln   True   -      66s
default   │   └─ Pod/silly-demo-5b764b57cc-rsw86   True   -      66s
```

Difference in behaviour can be seen if we apply same deployment with different image. Pods will be updated in rolling update manner, one by one.

```
kubectl apply --filename deployment/image.yaml \
```

```
&& watch kubectl tree deployment silly-demo
```

```
NAMESPACE NAME                                READY REASON AGE
default   Deployment/silly-demo                    -      -      4m59s
default   └─ ReplicaSet/silly-demo-5b764b57cc      -      -      4m59s
default   │   └─ Pod/silly-demo-5b764b57cc-gxtln   True   -      4m59s
default   │   └─ Pod/silly-demo-5b764b57cc-r8rc9   True   -      3s
default   │   └─ Pod/silly-demo-5b764b57cc-rsw86   True   -      4m59s
default   │   └─ Pod/silly-demo-5b764b57cc-zzn2z   True   -      3s
default   └─ ReplicaSet/silly-demo-75fd5bcc7c      -      -      3s
default   │   └─ Pod/silly-demo-75fd5bcc7c-brfv7   False  ContainersNotReady 2s
default   │   └─ Pod/silly-demo-75fd5bcc7c-gwtbd   False  ContainersNotReady 3s
```

NAMESPACE	NAME	READY	REASON	AGE
default	Deployment/silly-demo	-		5m32s
default	ReplicaSet/silly-demo-5b764b57cc	-		5m32s
default	ReplicaSet/silly-demo-75fd5bcc7c	-		36s
default	Pod/silly-demo-75fd5bcc7c-brfv7	True		35s
default	Pod/silly-demo-75fd5bcc7c-gwtbd	True		36s
default	Pod/silly-demo-75fd5bcc7c-qf98d	True		21s
default	Pod/silly-demo-75fd5bcc7c-tl7gg	True		9s
default	Pod/silly-demo-75fd5bcc7c-wdgfs	True		21s

After a bit of a time, all pods from new replicaset got rolled out.

Deployment can manage PV, but you need to create PVC at first, then connect it to deployment and attach mount point on each container:

```
cat deployment/volume.yaml
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: silly-claim
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
---
```

```
  readinessProbe:
    httpGet:
      path: /
      port: 8080
  resources:
    limits:
      cpu: 250m
      memory: 256Mi
    requests:
      cpu: 125m
      memory: 128Mi
  volumeMounts:
    - mountPath: /cache
      name: silly-cache
  volumes:
    - name: silly-cache
      persistentVolumeClaim:
        claimName: silly-claim
```


apply

```
kubectl apply --filename deployment/volume.yaml
```

```
kubectl get pods,persistentvolumes
```

```
controlplane $ kubectl get pods,persistentvolumes
```

NAME	READY	STATUS	RESTARTS	AGE
pod/silly-demo-75fd5bcc7c-brfv7	1/1	Running	0	26m
pod/silly-demo-75fd5bcc7c-gwrtbd	1/1	Running	0	26m
pod/silly-demo-75fd5bcc7c-qf98d	1/1	Running	0	25m
pod/silly-demo-75fd5bcc7c-wdgfs	1/1	Running	0	25m
pod/silly-demo-9dc9db44c-gbccr	1/1	Running	0	11s
pod/silly-demo-9dc9db44c-lnq2g	1/1	Running	0	11s

NAME	ATTRIBUTES	CLASS	REASON	AGE	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	VOLUME
persistentvolume/pvc-abf40891-c705-4f0c-b135-3f3f351d3254				7s	1Gi	RWO	Delete	Bound	default/silly-claim	local-path	<unse

As you can see, 1 PV is created for all Pods. A Deployment can manage Pods with PVCs, but it doesn't provide guarantees about Pod identity or stable network identifiers. Delete deployment.

```
kubectl delete --filename deployment/volume.yaml
```

Kubernetes Workload API: StatefulSet

StatefulSet Use Case: Running databases or applications needing persistent storage or unique configurations. Unlike Deployments, StatefulSets automatically manage PersistentVolumeClaims for each Pod.

apply and check

```
kubectl apply --filename statefulset/base.yaml
```

```
watch kubectl tree statefulset silly-demo
```

NAMESPACE	NAME	READY	REASON	AGE
default	StatefulSet/silly-demo	-		26s
default	ControllerRevision/silly-demo-6848df9f6f	-		26s
default	Pod/silly-demo-0	True		26s
default	Pod/silly-demo-1	False	ContainersNotReady	6s

Each created pod is numbered from 0 up. Check other specs:

```
kubectl get pods,persistentvolumes
```

```
controlplane $ kubectl get pods,persistentvolumes
NAME          READY  STATUS   RESTARTS  AGE
pod/silly-demo-0  1/1    Running    0          4m3s
pod/silly-demo-1  1/1    Running    0          3m43s

NAME          CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS  CLAIM
persistentvolume/pvc-a9d2607b-4b76-4260-bef5-f3ae0281f8a9  1Gi      RWO           Delete         Bound    default/silly-c
nset>
persistentvolume/pvc-b6b41cf7-b2fb-41bd-9b01-b71c6e6e8dc1  1Gi      RWO           Delete         Bound    default/silly-c
nset>
```

As expected, each pod has its own volume. Notice that StatefulSet doesn't make use of ReplicaSet. Each pod is created separately.

Change number of pods and lets see what will happen:

```
kubectl apply --filename statefulset/replicas.yaml \
&& watch kubectl tree statefulset silly-demo
```

NAMESPACE	NAME	READY	REASON	AGE
default	StatefulSet/silly-demo	-		20m
default	ControllerRevision/silly-demo-6848df9f6f	-		20m
default	Pod/silly-demo-0	True		20m
default	Pod/silly-demo-1	True		20m
default	Pod/silly-demo-2	True		24s
default	Pod/silly-demo-3	-		4s

New pods are created, one after another with according numeration. After deletion of statefullset, each pod and related PV will be deleted. From pods with biggest number, to the lowest.

```
kubectl delete --filename statefulset/replicas.yaml
```

Kubernetes Workload API: DaemonSet

DaemonSet use case: Running node-specific services like logging, monitoring agents, or network daemons.

apply and check manifest:

```
kubectl apply --filename daemonset/base.yaml  
cat daemonset/base.yaml
```

An important point about DaemonSet is that its manifest does not define the number of replicas.

```
controlplane $ kubectl apply --filename daemonset/base.yaml  
daemonset.apps/silly-demo created  
controlplane $ cat daemonset/base.yaml  
apiVersion: apps/v1  
kind: DaemonSet  
metadata:  
  name: silly-demo  
  labels:  
    app.kubernetes.io/name: silly-demo  
spec:  
  selector:  
    matchLabels:  
      app.kubernetes.io/name: silly-demo  
  template:  
    metadata:  
      labels:  
        app.kubernetes.io/name: silly-demo  
    spec:  
      containers:  
        - image: ghcr.io/vfarcic/silly-demo:1.4.116  
          name: silly-demo  
      ports:
```

```
kubectl tree daemonset silly-demo
```

```
controlplane $ kubectl tree daemonset silly-demo
NAMESPACE NAME READY REASON AGE
default DaemonSet/silly-demo - 5m23s
default ControllerRevision/silly-demo-6cb4b99cf8 - 5m23s
default Pod/silly-demo-7z74q True 5m23s
default Pod/silly-demo-d6hs8 True 5m23s
controlplane $
```

Even without defining replica count, DaemonSet creates Pods. What determines their number?
The number of cluster nodes.

```
kubectl get nodes
```

```
controlplane $ kubectl get nodes
NAME STATUS ROLES AGE VERSION
controlplane Ready control-plane 18d v1.31.0
node01 Ready <none> 18d v1.31.0
```

The number of nodes in the cluster is directly correlated with the number of pods created by a DaemonSet. For each node in the cluster, the DaemonSet ensures that exactly one pod is scheduled and running. As nodes are added or removed, the number of pods managed by the DaemonSet automatically adjusts to match the updated node count.

Delete workload.

```
kubectl delete --filename daemonset/base.yaml
```

Kubernetes Workload API: CronJob

Job use case: Periodic tasks like backups, reports, or maintenance scripts.

Check manifest. CronJob manifest is same as Job, with little adjustment. There is an CronJob defined.

Job will be ran every 1 minute:

```
cat cronjob/base.yaml
```

```
controlplane $ cat cronjob/base.yaml
apiVersion: batch/v1
kind: CronJob
metadata:
  name: silly-demo
  labels:
    app.kubernetes.io/name: silly-demo
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      template:
        metadata:
          labels:
            app.kubernetes.io/name: silly-demo
        spec:
          restartPolicy: OnFailure
          containers:
            - image: cgr.dev/chainguard/bash
              name: silly-demo
              command: ["echo", "What is this?"]
```

```
apply
```

```
kubectl apply --filename cronjob/base.yaml
```

```
watch kubectl get pods
```

Pods are created every minute:

```
Every 2.0s: kubectl get pods
NAME                                READY  STATUS   RESTARTS  AGE
silly-demo-28917341-58gp9           0/1    Completed 0          2m3s
silly-demo-28917342-8bxmf           0/1    Completed 0          63s
silly-demo-28917343-kv8cn           0/1    Completed 0          3s
```

common troubleshooting

1) ReplicaSet Not Scaling Pods

Cause: Misconfigured ReplicaSet spec or insufficient resources.

Solution: Check `kubectl get replicaset` for replicas and resource availability. Review events with `kubectl describe replicaset`.

2) Deployment Not Updating

Cause: Spec issues, image errors, or rollout failures.

Solution: Inspect Deployment with `kubectl describe`. Check image errors, rollout strategy, or undo with `kubectl rollout undo`.

3) StatefulSet Pods Stuck in Pending

Cause: Unbound volumes or insufficient resources.


Solution: Check PVCs with `kubectl get pvc`, and node resources with `kubectl describe nodes`. Verify storage and events.

4) Check my Kubernetes Troubleshooting series:

Michael Robotics

Hi, I'm Michal. I'm a Robotics Engineer and DevOps enthusiast. My mission is to create skill-learning platform that combats skill information overload by adhering to the set of principles: simplify, prioritize, and execute.

 <https://github.com/MichaelRobotics>




Learn more about Kubernetes

Check Kubernetes and piyushsachdeva - great docs!

Setup a Multi Node Kubernetes Cluster

kubeadm is a tool to bootstrap the Kubernetes cluster

 <https://github.com/piyushsachdeva/CKA-2024/tree/main/Resources/Day27>



Kubernetes Documentation

This section lists the different ways to set up and run Kubernetes

 <https://kubernetes.io/docs/setup/>



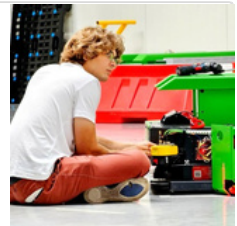
Share, comment, DM and check GitHub for scripts & playbooks created to automate process.

Check my GitHub

Michael Robotics

Hi, I'm Michal. I'm a Robotics Engineer and DevOps enthusiast. My mission is to create skill-learning platform that combats skill information overload by adhering to the set of principles: simplify, prioritize, and execute.

<https://github.com/MichaelRobotics>



PS.

If you need a playbook or bash script to manage KVM on a specific Linux distribution, feel free to ask me in the comments or send a direct message!